

Optimal Control and Estimation

Heng Yang

2023-06-29

Contents

Preface	5
1 The Optimal Control Formulation	7
1.1 The Basic Problem	7
1.2 Dynamic Programming and Principle of Optimality	9
2 Exact Dynamic Programming	13
2.1 Linear Quadratic Regulator	13
3 Approximate Dynamic Programming	15
4 Stability Analysis	17
5 Output Feedback	19
5.1 State Observer	19
6 Adaptive Control	29
6.1 Model-Reference Adaptive Control	29
6.2 Certainty-Equivalent Adaptive Control	39
A Linear System Theory	41
B Convex Analysis and Optimization	43
C The Kalman-Yakubovich Lemma	45
D Feedback Linearization	47
E Sliding Control	49

Preface

This is the textbook for Harvard ES/AM 158: Introduction to Optimal Control and Estimation. Information about the offerings of the class is listed below.

2023 Fall

Time: Mon/Wed 2:15 - 3:30pm

Location: Science and Engineering Complex, Room TBD

Instructor: Heng Yang

Teaching Fellow: Weiyu Li

Syllabus

Acknowledgment

Chapter 1

The Optimal Control Formulation

1.1 The Basic Problem

Consider a discrete-time dynamical system

$$x_{k+1} = f_k(x_k, u_k, w_k), \quad k = 0, 1, \dots, N-1 \quad (1.1)$$

where

- $x_k \in \mathbb{X} \subseteq \mathbb{R}^n$ is the *state* of the system,
- $u_k \in \mathbb{U} \subseteq \mathbb{R}^m$ is the *control* we wish to design,
- $w_k \in \mathbb{W} \subseteq \mathbb{R}^p$ a random *disturbance* or noise (e.g., due to unmodelled dynamics) which is described by a probability distribution $P_k(\cdot \mid x_k, u_k)$ that may depend on x_k and u_k but not on prior disturbances w_0, \dots, w_{k-1} ,
- k indexes the discrete time,
- N denotes the horizon,
- f_k models the transition function of the system (typically $f_k \equiv f$ is time-invariant, especially for robotics systems; we use f_k here to keep full generality).

Remark (Deterministic v.s. Stochastic). When $w_k \equiv 0$ for all k , we say the system (1.1) is *deterministic*; otherwise we say the system is *stochastic*. In the following we will deal with the stochastic case, but most of the methodology should carry over to the deterministic setup.

We consider the class of *controllers* (also called *policies*) that consist of a sequence of functions

$$\pi = \{\mu_0, \dots, \mu_{N-1}\},$$

where $\mu_k(x_k) \in \mathbb{U}$ for all x_k , i.e., μ_k is a *feedback* controller that maps the state to an admissible control. Given an initial state x_0 and an admissible policy π , the state *trajectory* of the system is a sequence of random variables that evolve according to

$$x_{k+1} = f_k(x_k, \mu_k(x_k), w_k), \quad k = 0, \dots, N-1 \quad (1.2)$$

where the randomness comes from the disturbance w_k .

We assume the state-control trajectory $\{u_k\}_{k=0}^{N-1}$ and $\{x_k\}_{k=0}^N$ induce an *additive cost*

$$g_N(x_N) + \sum_{k=0}^{N-1} g_k(x_k, u_k) \quad (1.3)$$

where $g_k, k = 0, \dots, N$ are some user-designed functions.

With (1.2) and (1.3), for any admissible policy π , we denote its induced *expected cost* with initial state x_0 as

$$J_\pi(x_0) = \mathbb{E} \left\{ g_N(x_N) + \sum_{k=0}^{N-1} g_k(x_k, \mu_k(x_k)) \right\}, \quad (1.4)$$

where the expectation is taken over the randomness of w_k .

Definition 1.1 (Discrete-time, Finite-horizon Optimal Control). Find the best admissible controller that minimizes the expected cost in (1.4)

$$\pi^* \in \arg \min_{\pi \in \Pi} J_\pi(x_0), \quad (1.5)$$

where Π is the set of all admissible controllers. The cost attained by the optimal controller, i.e., $J^* = J_{\pi^*}(x_0)$ is called the optimal *cost-to-go*, or the optimal *value function*.

Remark (Open-loop v.s. Closed-loop). An important feature of the basic problem in Definition 1.1 is that the problem seeks *feedback policies*, instead of numerical values of the controls, i.e., $u_k = \mu_k(x_k)$ is in general a function of the state x_k . In other words, the controls are executed sequentially, one at a time after observing the state at each time. This is called closed-loop control, and is in general better than open-loop control

$$\min_{u_0, \dots, u_{N-1}} \mathbb{E} \left\{ g_N(x_N) + \sum_{k=0}^{N-1} g_k(x_k, u_k) \right\}$$

where all the controls are planned at $k = 0$. Intuitively, a closed-loop policy is able to utilize the extra information received at each timestep (i.e., it observes x_{k+1} and hence also observes the disturbance w_k) to obtain a lower cost than an open-loop controller. Example 1.2.1 in (?) gives a concrete application where a closed-loop policy attains a lower cost than an open-loop policy.

In deterministic control (i.e., when $w_k \equiv 0, \forall k$), however, a closed-loop policy has no advantage over an open-loop controller. This is obvious because at $k = 0$, even the open-loop controller predicts perfectly the consequences of all its actions and there is no extra information to be observed at later time steps. In fact, even in stochastic problems, a closed-loop policy may not be advantageous, see Exercise 1.27 in (?).

1.2 Dynamic Programming and Principle of Optimality

We now introduce a general and powerful algorithm, namely *dynamic programming* (DP), for solving the optimal control problem 1.1. The DP algorithm builds upon a quite simple intuition called the *Bellman principle of optimality*.

Theorem 1.1 (Bellman Principle of Optimality). *Let $\pi^* = \{\mu_0^*, \mu_1^*, \dots, \mu_{N-1}^*\}$ be an optimal policy for the optimal control problem 1.1. Assume that when using π^* , a given state x_i occurs at timestep i with positive probability (i.e., x_i is reachable at time i).*

Now consider the following subproblem where we are at x_i at time i and wish to minimize the cost-to-go from time i to time N

$$\min_{\mu_i, \dots, \mu_{N-1}} \mathbb{E} \left\{ g_N(x_N) + \sum_{k=i}^{N-1} g_k(x_k, \mu_k(x_k)) \right\}.$$

Then the truncated policy $\{\mu_i^, \mu_{i+1}^*, \dots, \mu_{N-1}^*\}$ must be optimal for the subproblem.*

Theorem 1.1 can be proved intuitively by contradiction: if the truncated policy $\{\mu_i^*, \mu_{i+1}^*, \dots, \mu_{N-1}^*\}$ is not optimal for the subproblem, say there exists a different policy $\{\mu'_i, \mu'_{i+1}, \dots, \mu'_{N-1}\}$ that attains a lower cost for the subproblem starting at x_i at time i . Then the combined policy $\{\mu_0^*, \dots, \mu_{i-1}^*, \mu'_i, \dots, \mu'_{N-1}\}$ must attain a lower cost for the original optimal control problem 1.1 due to the additive cost structure, contradicting the optimality of π^* .

The Bellman principle of optimality is more than just a principle, it is also an algorithm. It suggests that, to build an optimal policy, one can start by solving the last-stage subproblem to obtain $\{\mu_{N-1}^*\}$, and then proceed to solve the subproblem containing the last two stages to obtain $\{\mu_{N-2}^*, \mu_{N-1}^*\}$. The recursion continues until optimal policies at all stages are computed. The following theorem formalizes this concept.

Theorem 1.2 (Dynamic Programming). *The optimal value function $J^*(x_0)$ of the optimal control problem 1.1 (starting from any given initial condition x_0) is*

equal to $J_0(x_0)$, which can be computed backwards and recursively as

$$J_N(X_N) = g_N(x_N) \quad (1.6)$$

$$J_k(x_k) = \min_{u_k \in \mathbb{U}} \mathbb{E}_{w_k \sim P_k(\cdot | x_k, u_k)} \{g_k(x_k, u_k) + J_{k+1}(f_k(x_k, u_k, w_k))\}, \quad k = N-1, \dots, 1, 0. \quad (1.7)$$

Moreover, if $u_k^* = \mu_k^*(x_k)$ is a minimizer of (1.7) for every x_k , then the policy $\pi^* = \{\mu_0^*, \dots, \mu_{N-1}^*\}$ is optimal.

Proof. For any admissible policy $\pi = \{\mu_0, \dots, \mu_{N-1}\}$, denote $\pi^k = \{\mu_k, \dots, \mu_{N-1}\}$ the last- $(N-k)$ -stage truncated policy. Consider the subproblem consisting of the last $N-k$ stages starting from x_k , and let $J_k^*(x_k)$ be its optimal cost-to-go. Mathematically, this is

$$J_k^*(x_k) = \min_{\pi^k} \mathbb{E}_{w_k, \dots, w_{N-1}} \left\{ g_N(x_N) + \sum_{i=k}^{N-1} g_i(x_i, \mu_i(x_i)) \right\}, \quad k = 0, 1, \dots, N-1. \quad (1.8)$$

We define $J_N^*(x_N) = g(x_N)$ for $k = N$.

Our goal is to prove the $J_k(x_k)$ computed by dynamic programming from (1.7) is equal to $J_k^*(x_k)$ for all $k = 0, \dots, N$. We will prove this by induction.

Firstly, we already have $J_N^*(x_N) = J_N(x_N) = g(x_N)$, so $k = N$ holds automatically.

Now we assume $J_{k+1}^*(x_{k+1}) = J_{k+1}(x_{k+1})$ for all x_{k+1} , and we wish to induce

$J_k^*(x_k) = J_k(x_k)$. To show this, we write

$$J_k^*(x_k) = \min_{\pi^k} \mathbb{E}_{w_k, \dots, w_{N-1}} \left\{ g_N(x_N) + \sum_{i=k}^{N-1} g_i(x_i, \mu_i(x_i)) \right\} \quad (1.9)$$

$$= \min_{\mu_k, \pi^{k+1}} \mathbb{E}_{w_k, \dots, w_{N-1}} \left\{ g_k(x_k, \mu_k(x_k)) + g_N(x_N) + \sum_{i=k+1}^{N-1} g_i(x_i, \mu_i(x_i)) \right\} \quad (1.10)$$

$$= \min_{\mu_k} \left[\min_{\pi^{k+1}} \mathbb{E}_{w_k, \dots, w_{N-1}} \left\{ g_k(x_k, \mu_k(x_k)) + g_N(x_N) + \sum_{i=k+1}^{N-1} g_i(x_i, \mu_i(x_i)) \right\} \right] \quad (1.11)$$

$$= \min_{\mu_k} \mathbb{E}_{w_k} \left\{ g_k(x_k, \mu_k(x_k)) + \min_{\pi^{k+1}} \left[\mathbb{E}_{w_{k+1}, \dots, w_{N-1}} \left\{ g_N(x_N) + \sum_{i=k+1}^{N-1} g_i(x_i, \mu_i(x_i)) \right\} \right] \right\} \quad (1.12)$$

$$= \min_{\mu_k} \mathbb{E}_{w_k} \{ g_k(x_k, \mu_k(x_k)) + J_{k+1}^*(f_k(x_k, \mu_k(x_k), w_k)) \} \quad (1.13)$$

$$= \min_{\mu_k} \mathbb{E}_{w_k} \{ g_k(x_k, \mu_k(x_k)) + J_{k+1}(f_k(x_k, \mu_k(x_k), w_k)) \} \quad (1.14)$$

$$= \min_{u_k \in \mathbb{U}} \mathbb{E}_{w_k} \{ g_k(x_k, \mu_k(x_k)) + J_{k+1}(f_k(x_k, \mu_k(x_k), w_k)) \} \quad (1.15)$$

$$= J_k(x_k), \quad (1.16)$$

where (1.9) follows from definition (1.8); (1.10) expands $\pi^k = \{\mu_k, \pi^{k+1}\}$ and $\sum_{i=k}^{N-1} g_i = g_k + \sum_{i=k+1}^{N-1} g_i$; (1.11) writes the joint minimization over μ_k and π^{k+1} as equivalently first minimizing over π^{k+1} and then minimizing over μ_k ; (1.12) is the key step and holds because g_k and w_k depend only on μ_k but not on π^{k+1} ; (1.13) follows again from definition (1.8) with k replaced by $k+1$; (1.14) results from the induction assumption; (1.15) clearly holds because any $\mu_k(x_k)$ belongs to \mathbb{U} and any element in \mathbb{U} can be chosen by a feedback controller μ_k ; and lastly (1.16) follows from the dynamic programming algorithm (1.7).

By induction, this shows that $J_k^*(x_k) = J_k(x_k)$ for all $k = 0, \dots, N$. \square

The careful reader, especially from a robotics background, may soon become disappointed when seeing the DP algorithm (1.7) because it is rather conceptual than practical. To see this, we only need to run DP for $k = N-1$:

$$J_{N-1}(x_{N-1}) = \min_{u_{N-1} \in \mathbb{U}} \mathbb{E}_{w_{N-1}} \{ g_{N-1}(x_{N-1}, u_{N-1}) + J_N(f_{N-1}(x_{N-1}, u_{N-1}, w_{N-1})) \}. \quad (1.17)$$

Two challenges immediately show up:

- How to perform the minimization over u_{N-1} when \mathbb{U} is a continuous constraint set? Even if we assume g_{N-1} is convex¹ in u_{N-1} , J_N is convex in

¹You may want to read Appendix B if this is your first time seeing “convex” things.

x_N , and the dynamics f_{N-1} is also convex in u_{N-1} (so that the optimization (1.17) is convex), we may be able to solve the minimization *numerically* for each x_{N-1} using a convex optimization solver, but rarely will we be able to find an analytical policy μ_{N-1}^* such that $u_{N-1}^* = \mu_{N-1}^*(x_{N-1})$ for every x_{N-1} (i.e., the optimal policy μ_{N-1}^* is implicit but not explicit).

- Suppose we can find an analytical optimal policy μ_{N-1}^* , say $\mu_{N-1}^* = Kx_{N-1}$ a linear policy, how will plugging μ_{N-1}^* into (1.17) affect the complexity of $J_{N-1}(x_{N-1})$? One can see that even if μ_{N-1}^* is linear in x_{N-1} , J_{N-1} may be highly nonlinear in x_{N-1} due to the composition with g_{N-1} , f_{N-1} and J_N . If $J_{N-1}(x_{N-1})$ becomes too complex, then clearly it becomes more challenging to perform (1.17) for the next step $k = N - 2$.

Due to these challenges, only in a very limited amount of cases will we be able to perform *exact dynamic programming*. For example, when the state space \mathbb{X} and control space \mathbb{U} are discrete, we can design efficient algorithms for exact DP. For another example, when the dynamics f_k is linear and the cost g_k is quadratic, we will also be able to compute $J_k(x_k)$ in closed form (though this sounds a bit surprising!). We will study these problems in more details in Chapter 2.

For general optimal control problems with continuous state space and control space (and most problems we care about in robotics), unfortunately, we will have to resort to *approximate dynamic programming*, basically variations of the DP algorithm (1.7) where approximate value functions $J_k(x_k)$ and/or control policies $\mu_k(x_k)$ are used (e.g., with neural networks and machine learning).² We will introduce several popular approximation schemes in Chapter 3. We will see that, although exact DP is not possible anymore, the Bellman principle of optimality still remains one of the most important guidelines for designing approximation algorithms. Efficient algorithms for approximate dynamic programming, preferably with performance guarantees, still remain an active area of research.

²Another possible solution is to discretize continuous states and controls. However, when the dimension of state and control is high, discretization becomes too expensive in terms of memory and computational complexity.

Chapter 2

Exact Dynamic Programming

2.1 Linear Quadratic Regulator

Chapter 3

Approximate Dynamic Programming

Chapter 4

Stability Analysis

Lemma 4.1 (Barbalat's Lemma). *Let $f(t)$ be differentiable, if*

- *$\lim_{t \rightarrow \infty} f(t)$ is finite, and*
- *$\dot{f}(t)$ is uniformly continuous,¹*

then

$$\lim_{t \rightarrow \infty} \dot{f}(t) = 0.$$

Theorem 4.1 (Barbalat's Stability Certificate). *If a scalar function $V(x, t)$ satisfies*

- *$V(x, t)$ is lower bounded,*
- *$\dot{V}(x, t)$ is negative semidefinite*
- *$\dot{V}(x, t)$ is uniformly continuous*

then $\dot{V}(x, t) \rightarrow 0$ as $t \rightarrow \infty$.

Proof. $V(x, t)$ is lower bounded and \dot{V} is negative semidefinite implies the limit of V as $t \rightarrow \infty$ is finite (note that $V(x, t) \leq V(x(0), 0)$). Then the theorem clearly follows from Barbalat's Lemma 4.1. \square

¹A sufficient condition for this to hold is that \ddot{f} exists and is bounded.

Chapter 5

Output Feedback

Consider a continuous-time dynamical system

$$\begin{aligned}\dot{x} &= f(x, u) \\ y &= h(x, u)\end{aligned}\tag{5.1}$$

where $x(t) \in \mathbb{X} \subseteq \mathbb{R}^n$ the state of the system, $u(t) \in \mathbb{U} \subseteq \mathbb{R}^m$ the control (or input), $y(t) \in \mathbb{Y} \subseteq \mathbb{R}^d$ the output (i.e., measurement) of the state and control, and f, g the evolution and measurement functions (which are sufficiently smooth).

5.1 State Observer

For the system (5.1), let us denote

- $X(x_0, t_0; t; u)$ the solution at time t with input u and initial condition x_0 at time t_0 ; when $t_0 = 0$, we write $X(x_0; t; u)$
- $Y(x_0, t_0; t; u)$ the output at time t with input u and initial condition x_0 at time t_0 , i.e., $Y(x_0, t_0; t; u) = h(X(x_0, t_0; t; u), u(t))$; when $t_0 = 0$, we write $y_{x_0, u}(t)$;
- \mathcal{X}_0 a subset of \mathbb{X} containing the initial conditions we consider; for any $x_0 \in \mathcal{X}_0$, we write $\sigma_{\mathcal{X}}^+(x_0; u)$ the maximal time of existence of $X(x_0, \cdot; t; u)$ in a set \mathcal{X}
- \mathcal{U} the set of all sufficiently many times differentiable inputs $u : [0, +\infty) \rightarrow \mathbb{U}$.

The problem of state observation is to produce an estimated state $\hat{x}(t)$ of the true state $X(x_0, t_0; t; u)$ based on knowledge about the system (5.1) and information about the history of inputs $u_{[0, t]}$ and outputs $y_{[0, t]}$, so that $\hat{x}(t)$ asymptotically converges to $X(x_0, t_0; t; u)$, for any initial condition $x_0 \in \mathcal{X}_0$ and any input $u \in \mathcal{U}$.

There are multiple ways for solving the problem of state observation (see e.g., (?), (?)). Here we are particularly interested in the approach using a *state observer*, i.e., a dynamical system whose *internal state* evolves according to the history of inputs and outputs, from which a state estimation can be reconstructed that guarantees asymptotic convergence to the true state. We formalize this concept below.

Definition 5.1 (State Observer). A state observer for system (5.1) is a couple $(\mathcal{F}, \mathcal{T})$ such that

1. $\mathcal{F} : \mathbb{R}^l \times \mathbb{R}^m \times \mathbb{R}^d \rightarrow \mathbb{R}^l$ is continuous
2. \mathcal{T} is a family of continuous functions indexed by $u \in \mathcal{U}$ where each $\mathcal{T}_u : \mathbb{R}^l \times [0, +\infty) \rightarrow \mathbb{R}^n$ respects the causality condition

$$\forall \tilde{u} : [0, +\infty) \rightarrow \mathbb{R}^m, \forall t \in [0, +\infty), u_{[0,t]} = \tilde{u}_{[0,t]} \Rightarrow \mathcal{F}_u(\cdot, t) = \mathcal{F}_{\tilde{u}}(\cdot, t).$$

3. For any $u \in \mathcal{U}$, any $z_0 \in \mathbb{R}^l$, and any $x_0 \in \mathcal{X}_0$ such that $\sigma_{\mathbb{X}}^+(x_0; u) = +\infty$, any solution $Z(z_0; t; u, y_{x_0, u})$ ¹ to

$$\dot{z} = \mathcal{F}(z, u, y_{x_0, u}) \quad (5.2)$$

initialized at z_0 at time 0 with input u and $y_{x_0, u}$ exists on $[0, +\infty)$ and satisfies

$$\lim_{t \rightarrow \infty} \|\hat{X}(x_0, z_0; t; u) - X(x_0; t; u)\| = 0, \quad (5.3)$$

with

$$\hat{X}(x_0, z_0; t; u) = \mathcal{T}_u(Z(z_0; t; u, y_{x_0, u}), t). \quad (5.4)$$

In words, (i) the state observer maintains an internal state (or latent state) $z \in \mathbb{R}^l$ that evolves according to the latent dynamics \mathcal{F} in (5.2), where u and $y_{x_0, u}$ are inputs; (ii) an estimated state can be reconstructed from the internal state using \mathcal{T}_u as in (5.4); and (iii) the error between the estimated state and the true state (defined by a proper distance function $\|\cdot\|$ on \mathbb{X}) converges to zero.

If \mathcal{T}_u is the same for any $u \in \mathcal{U}$ and is also time independent, then we say \mathcal{T} is *stationary*.² In this case, we can simply write the observer (5.2) and (5.4) as

$$\begin{aligned} \dot{z} &= \mathcal{F}(z, u, y) \\ \hat{x} &= \mathcal{T}(z). \end{aligned} \quad (5.5)$$

¹We say “any solution” because there may be several solutions to the observer (5.2) due to \mathcal{F} only being continuous. This is not a problem as long as any such solution satisfies the required convergence property.

²The time dependence of \mathcal{T}_u enables us to cover the case where the knowledge of the u and $y_{x_0, u}$ is used to construct the estimate from the observer state. In particular, using the output sometimes can reduce the dimension of the observer state (and thus alleviate the computations), thus obtaining a reduced-order observer. For example, see (?) and (?).

If \hat{x} can be read off directly from z , then we say the observer (5.5) is *in the given coordinates*. A special case of this is when $\hat{x} = z$, i.e., the internal state of the observer is the same as the system state.

5.1.1 General Design Strategy

Theorem 5.1 (Meta Observer). *Let $F : \mathbb{R}^p \times \mathbb{R}^m \times \mathbb{R}^d \rightarrow \mathbb{R}^p$, $H : \mathbb{R}^p \times \mathbb{R}^m \rightarrow \mathbb{R}^d$ and $\mathcal{F} : \mathbb{R}^p \times \mathbb{R}^m \times \mathbb{R}^d \rightarrow \mathbb{R}^p$ be continuous functions such that*

$$\dot{\hat{\xi}} = \mathcal{F}(\hat{\xi}, u, \tilde{y}) \quad (5.6)$$

is an observer for

$$\dot{\xi} = F(\xi, u, H(\xi, u)), \quad \tilde{y} = H(\xi, u), \quad (5.7)$$

i.e., for any $\xi_0, \hat{\xi}_0 \in \mathbb{R}^p$ and any $u \in \mathcal{U}$, the solution of the observer (5.6), denoted by $\hat{\Xi}(\hat{\xi}_0; t; u; \tilde{y}_{\xi_0, u})$, and the solution of the true system (5.7), denoted by $\Xi(\xi_0; t; u)$, satisfy

$$\lim_{t \rightarrow \infty} \|\hat{\Xi}(\hat{\xi}_0; t; u; \tilde{y}_{\xi_0, u}) - \Xi(\xi_0; t; u)\| = 0. \quad (5.8)$$

Note that the observer (5.6) is stationary and in the given coordinates for system (5.7). Indeed the internal state of the observer is the same as the system state.

Now suppose for any $u \in \mathcal{U}$, there exists a continuous function (i.e., coordinate transformation) $T_u : \mathbb{R}^n \times \mathbb{R} \rightarrow \mathbb{R}^p$ and a subset \mathcal{X} of \mathbb{X} such that

1. *For any $x_0 \in \mathcal{X}_0$ such that $\sigma_{\mathbb{X}}^+(x_0; u) = +\infty$, $X(x_0; \cdot; u)$ remains in \mathcal{X}*
2. *There exists a concave \mathcal{K}^3 function ρ and a positive number \bar{t} such that*

$$\|x_a - x_b\| \leq \rho(|T_u(x_a, t) - T_u(x_b, t)|), \quad \forall x_a, x_b \in \mathcal{X}, t \geq \bar{t},$$

i.e., $x \mapsto T_u(x, t)$ becomes injective on \mathcal{X} ,⁴ uniformly in time and space, after a certain time \bar{t} .

3. *T_u transforms the system (5.1) into the system (5.7), i.e., for all $x \in \mathcal{X}$ and all $t \geq 0$, we have*

$$L_{(f,1)}T_u(x, t) = F(T_u(x, t), u, h(x, u)), \quad h(x, u) = H(T_u(x, t), u), \quad (5.9)$$

where $L_{(f,1)}T_u(x, t)$ is the Lie derivative of T_u along the vector field $(f, 1)$

$$L_{(f,1)}T_u(x, t) = \lim_{\tau \rightarrow 0} \frac{T_u(X(x, t; t + \tau; u), t + \tau) - T_u(x, t)}{\tau}.$$

³A function $\rho : \mathbb{R}_+ \rightarrow \mathbb{R}_+$ is a \mathcal{K} function if $\rho(0) = 0$, ρ is continuous, and ρ is increasing.

⁴An injective function is a function f that maps distinct elements of its domain to distinct elements. That is, $f(x_a) = f(x_b)$ implies $x_a = x_b$, or equivalently, $x_a \neq x_b$ implies $f(x_a) \neq f(x_b)$.

4. T_u respects the causality condition

$$\forall \tilde{u} : [0, +\infty) \rightarrow \mathbb{R}^m, \forall t \in [0, +\infty), u_{[0,t]} = \tilde{u}_{[0,t]} \Rightarrow T_u(\cdot, t) = T_{\tilde{u}}(\cdot, t).$$

Then, for any $u \in \mathcal{U}$, there exists a function $\mathcal{T}_u : \mathbb{R}^p \times [0, +\infty) \rightarrow \mathcal{X}$ (satisfying the causality condition) such that for any $t \geq \bar{t}$, $\xi \mapsto \mathcal{T}_u(\xi, t)$ is uniformly continuous on \mathbb{R}^p and satisfies

$$\mathcal{T}_u(T_u(x, t), t) = x, \forall x \in \mathcal{X}.$$

Moreover, denoting \mathcal{T} the family of functions \mathcal{T}_u for $u \in \mathcal{U}$, the couple $(\mathcal{F}, \mathcal{T})$ is an observer for the system (5.1) initialized in \mathcal{X}_0 .

Proof. See Theorem 1.1 in (?). □

A simpler version of Theorem 5.1 where the coordinate transformation T_u is stationary and fixed for all u is stated below as a corollary.

Corollary 5.1 (Meta Observer with Fixed Transformation). *Let $F : \mathbb{R}^p \times \mathbb{R}^m \times \mathbb{R}^d \rightarrow \mathbb{R}^p$, $H : \mathbb{R}^p \times \mathbb{R}^m \rightarrow \mathbb{R}^d$ and $\mathcal{F} : \mathbb{R}^p \times \mathbb{R}^m \times \mathbb{R}^d \rightarrow \mathbb{R}^p$ be continuous functions such that (5.6) is an observer for (5.7).*

Suppose there exists a continuous coordinate transformation $T : \mathbb{R}^p \rightarrow \mathbb{R}^n$ and a compact subset Ω of \mathbb{R}^n such that

1. *For any $x_0 \in \mathcal{X}_0$ such that $\sigma_{\mathbb{X}}^+(x_0; u) = +\infty$, $X(x_0; \cdot; u)$ remains in Ω*
2. *$x \mapsto T(x)$ is injective on Ω*
3. *T transforms the system (5.1) into system (5.7)*

$$L_f T(x) = F(T(x), u, h(x, u)), \quad h(x, u) = H(T(x), u),$$

where $L_f T(x)$ is the Lie derivative of $T(x)$ along f

$$L_f T(x) = \lim_{\tau \rightarrow 0} \frac{T(X(x, t; t + \tau; u)) - T(x)}{\tau}.$$

Then, there exists a uniformly continuous function $\mathcal{T} : \mathbb{R}^p \rightarrow \mathbb{R}^n$ such that

$$\mathcal{T}(T(x)) = x, \quad \forall x \in \Omega,$$

and $(\mathcal{F}, \mathcal{T})$ is an observer for system (5.1) initialized in \mathcal{X}_0 .

Theorem 5.1 and Corollary 5.1 suggest the following general observer design strategy:

1. Find an injective coordinate transformation T_u (that may be time-varying and also dependent on u) that transforms the original system (5.1) with coordinate x into a new system (5.7) with coordinate ξ

2. Design an observer (5.6), $\hat{\xi}$, for the new system
3. Compute a left inverse, \mathcal{T}_u , of the transformation T_u to recover a state estimation \hat{x} of the original system.

The transformed systems (5.7) are typically referred to as *normal forms*, or in my opinion, *templates*.

Of course, the general design strategy is rather conceptual, and in order for it to be practical, we have to answer three questions.

- What templates do we have, what are their associated observers, and what are the conditions for the observers to be asymptotically converging?
- What kinds of (nonlinear) systems can be transformed into the templates, and how to perform the transformation?
- How to invert the coordinate transformation? Is it analytical or does it require numerical approximation?

In the following sections, we will study several representative normal forms and answer the above questions.

5.1.2 Luenberger Template

Consider an instance of the normal form (5.7) as follows:

$$\dot{\xi} = A\xi + B(u, y), \quad y = C\xi, \quad (5.10)$$

where A, C are constant matrices, and $B(u, y)$ can depend nonlinearly on u and y .

For this template, we have the well-known Luenberger observer.

Theorem 5.2 (Luenberger Observer). *If the pair (A, C) is detectable (see Theorem A.1), then there exists a matrix K such that $A - KC$ is Hurwitz and the system*

$$\dot{\hat{\xi}} = A\hat{\xi} + B(u, y) + K(y - C\hat{\xi}) \quad (5.11)$$

is an observer for (5.10).

Proof. TBD. □

The Luenberger observer is an elegant result in observer design (and even in control theory) that has far-reaching impact. In my opinion, the essence of observer design is twofold: (i) to simulate the dynamics when the state estimation is correct, and (ii) to correct the state estimation from observation when it is off. These two pieces of ideas are evident in (5.11): the observer is a copy of the original dynamics ($A\hat{\xi} + B(u, y)$) plus a feedback correction from the difference between the “imagined” observation $C\hat{\xi}$ and the true observation y .

You may think the Luenberger template is restricting because it requires the system to be linear (up to the only nonlinearity in $B(u, y)$). However, it turns out the Luenberger template is already quite useful, as I will show in the following pendulum example.

Example 5.1 (Luenberger Observer for A Simple Pendulum). Consider a simple pendulum dynamics model

$$x = \begin{bmatrix} \theta \\ \dot{\theta} \end{bmatrix}, \quad \dot{x} = \begin{bmatrix} \dot{\theta} \\ -\frac{1}{ml^2}(b\dot{\theta} + mgl \sin \theta) \end{bmatrix} + \begin{bmatrix} 0 \\ \frac{1}{ml^2} \end{bmatrix} u, \quad y = \theta, \quad (5.12)$$

where θ the angular position of the pendulum from the vertical line, $m > 0$ the mass of the pendulum, $l > 0$ the length, g the gravitational constant, $b > 0$ the damping coefficient, and u the control input (torque).

We assume we can only observe the angular position of the pendulum in (5.12), e.g., using a camera, but not the angular velocity. Our goal is to construct an observer that can provide a full state estimation.

We first note that the pendulum dynamics (5.12) can actually be written in the (linear) Luenberger template (5.10) as⁵

$$\begin{aligned} \dot{x} &= \underbrace{\begin{bmatrix} 0 & 1 \\ 0 & -\frac{b}{ml^2} \end{bmatrix}}_{=:A} x + \underbrace{\begin{bmatrix} 0 \\ \frac{u - mgl \sin \theta}{ml^2} \end{bmatrix}}_{=:B(u,y)} \\ y &= \underbrace{\begin{bmatrix} 1 & 0 \end{bmatrix}}_{=:C} x \end{aligned} \quad (5.13)$$

In order for us to use the Luenberger observer, we need to check if the pair (A, C) is detectable. We can easily find the eigenvalues and eigenvectors of A :

$$A \begin{bmatrix} 1 \\ 0 \end{bmatrix} = 0, \quad A \begin{bmatrix} -\frac{ml^2}{b} \\ 1 \end{bmatrix} = \begin{bmatrix} 1 \\ -\frac{b}{ml^2} \end{bmatrix} = -\frac{b}{ml^2} \begin{bmatrix} -\frac{ml^2}{b} \\ 1 \end{bmatrix}.$$

The first eigenvalue has real part equal to 0. However,

$$C \begin{bmatrix} 1 \\ 0 \end{bmatrix} = 1 \neq 0.$$

According to Theorem A.1, we conclude (A, C) is detectable.

Now we need to find K . An easy choice of K is

$$K = \begin{bmatrix} k \\ 0 \end{bmatrix}, \quad A - KC = \begin{bmatrix} -k & 1 \\ 0 & -\frac{b}{ml^2} \end{bmatrix}.$$

With $k > 0$, we know $A - KC$ is guaranteed to be Hurwitz (the two eigenvalues of $A - KC$ are $-k$ and $-b/ml^2$), and we have obtained an observer!

TBD: simulate the observer.

⁵I have to say I was a bit surprised when I arrived at this formulation.

5.1.3 State-affine Template

Consider an instance of the normal form (5.7) where the dynamics is linear in ξ , but the coefficients are time-varying and dependent on the input and output

$$\dot{\xi} = A(u, y)\xi + B(u, y), \quad y = C(u)\xi. \quad (5.14)$$

The difference between the state-affine template (5.14) and the Luenberger template (5.10) is that the linear matrices A, C are allowed to depend nonlinearly on the input (u, y) .

Kalman and Bucy originally proposed an observer for linear time-varying systems (?). The result is later extended by (?) and (?) to deal with coefficient matrices dependent on the control. The following theorem is a direct extension of the result from (?) and (?) by considering (u, y) as an augmented control input.

Before presenting the theorem, we need to introduce the following terminology.

Definition 5.2 (Linear Time-Varying System). For a linear time-varying system of the form

$$\dot{\chi} = A(\nu)\chi, \quad y = C(\nu)\chi, \quad (5.15)$$

with input ν and output y , we define

- the *transition matrix* Ψ_ν as the unique solution to

$$\Psi_\nu(t, t) = I, \quad \frac{\partial \Psi_\nu}{\partial \tau}(\tau, t) = A(\nu(\tau))\Psi_\nu(\tau, t).$$

Note that the transition matrix is used to express the solution to (5.15) because it satisfies

$$\chi(\chi_0, t_0; t; \nu) = \Psi_\nu(t, t_0)\chi_0.$$

- the *observability grammian* as

$$\Gamma_\nu(t_0, t_1) = \int_{t_0}^{t_1} \Psi_\nu(\tau, t_0)^T C(\nu(\tau))^T C(\nu(\tau)) \Psi_\nu(\tau, t_0) d\tau.$$

- the *backward observability grammian* as

$$\Gamma_\nu^b(t_0, t_1) = \int_{t_0}^{t_1} \Psi_\nu(\tau, t_1)^T C(\nu(\tau))^T C(\nu(\tau)) \Psi_\nu(\tau, t_1) d\tau.$$

We now introduce the Kalman-Bucy Observer for the state-affine template (5.14).

Theorem 5.3 (Kalman-Bucy Observer). *Let $y_{\xi_0, u}(t) = C(u(t))\Xi(\xi_0; t; u)$ be the output of system (5.14) at time t with initialization ξ_0 and control u . Suppose the control u satisfies*

- For any ξ_0 , $t \mapsto A(u(t), y_{\xi_0, u}(t))$ is bounded by A_{\max}
- For any ξ_0 , the augmented input $\nu = (u, y_{\xi_0, u})$ is regularly persistent for the dynamics

$$\dot{\chi} = A(\nu)\chi, \quad y = C(\nu)\chi \quad (5.16)$$

uniformly with respect to ξ_0 . That is, there exist strictly positive numbers t_0, \bar{t} , and α such that for any ξ_0 and any time $t \geq t_0 \geq \bar{t}$,

$$\Gamma_v^b(t - \bar{t}, t) \succeq \alpha I,$$

where Γ_v^b is the backward observability grammian associated with system (5.16).

Then, given any positive definite matrix P_0 , there exist $\alpha_1, \alpha_2 > 0$ such that for any $\lambda \geq 2A_{\max}$ and any $\xi_0 \in \mathbb{R}^p$, the matrix differential equation

$$\dot{P} = -\lambda P - A(u, y)^T P - PA(u, y) + C(u)^T C(u) \quad (5.17)$$

initialized at $P(0) = P_0$ admits a unique solution satisfying $P(t) = P(t)^T$ and

$$\alpha_2 I \succeq P(t) \succeq \alpha_1 I.$$

Moreover, the system

$$\dot{\hat{\xi}} = A(u, y)\hat{\xi} + B(u, y) + K(y - C(u)\hat{\xi}) \quad (5.18)$$

with a time-varying gain matrix

$$K = P^{-1}C(u)^T \quad (5.19)$$

is an observer for the state-affine system (5.14).

Let us work out an example of the Kalman-Bucy Observer for nonlinear systems.

Example 5.2 (Kalman-Bucy Observer for A Simple Pendulum). Let us reconsider the pendulum dynamics (5.12) but this time try to design a Kalman-Bucy observer.

We first write the pendulum dynamics in a new coordinate system so that it is in the state-affine normal form (5.14). We choose $\xi = [\mathfrak{s}, \mathfrak{c}, \dot{\theta}]^T$ with $\mathfrak{s} = \sin \theta$ and $\mathfrak{c} = \cos \theta$. Clearly, we will be able to observe $y = [\mathfrak{s}, \mathfrak{c}]^T$ in this new coordinate. The state-affine normal form of the pendulum dynamics reads

$$\begin{aligned} \dot{\xi} &= \begin{bmatrix} \mathfrak{c}\dot{\theta} \\ -\mathfrak{s}\dot{\theta} \\ -\frac{1}{ml^2}(b\dot{\theta} + mgl\mathfrak{s}) + \frac{1}{ml^2}u \end{bmatrix} = \underbrace{\begin{bmatrix} 0 & 0 & \mathfrak{c} \\ 0 & 0 & -\mathfrak{s} \\ 0 & 0 & -\frac{b}{ml^2} \end{bmatrix}}_{=:A(u,y)} \xi + \underbrace{\begin{bmatrix} 0 \\ 0 \\ \frac{u-mgl\mathfrak{s}}{ml^2} \end{bmatrix}}_{=:B(u,y)} \\ y &= \underbrace{\begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \end{bmatrix}}_{=:C(u)} \xi \end{aligned} \quad (5.20)$$

Note that $C(u)$ is in fact time-invariant, and $B(u, y)$ only depends on u ; but we adopt the same notation as the general state-affine template (5.14).

In order to use the Kalman-Bucy observer in Theorem 5.3, we need to verify the boundedness of $A(u, y)$, and the regular persistence of (5.16).

Boundedness of $A(u, y)$. We can easily show the boundedness of $A(u, y)$ by writing

$$\|A(u, y)\xi\| = \|\xi_3(\mathfrak{c} - \mathfrak{s} - b/ml^2)\| \leq |\xi_3| \sqrt{3} \sqrt{\mathfrak{c}^2 + \mathfrak{s}^2 + b^2/m^2 l^4} \leq \|\xi\| \sqrt{3 + 3b^2/m^2 l^4}.$$

Therefore, we can take $A_{\max} = \sqrt{3 + 3b^2/m^2 l^4}$. Does the A_{\max} in Theorem 5.3 refer to the bound in operator norm?

Regular persistence. We write the backward observability grammian of system (5.16)

$$\Gamma_{\nu}^b(t - \bar{t}, t) = \int_{t - \bar{t}}^t \Psi_{\nu}(\tau, t)^T C^T C \Psi_{\nu}(\tau, t) d\tau = \int_{t - \bar{t}}^t \Psi_{\nu}(\tau, t)^T \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix} \Psi_{\nu}(\tau, t) d\tau$$

Does the transition matrix $\Psi_{\nu}(\tau, t)$ have an analytical form?

5.1.4 Triangular Template

Chapter 6

Adaptive Control

6.1 Model-Reference Adaptive Control

Basic flow for designing an adaptive controller

1. Design a control law with variable parameters
2. Design an adaptation law for adjusting the control parameters
3. Analyze the convergence of the closed-loop system

The control law design at the first step typically requires the designer to know what a good controller is if the true parameters were actually known, e.g., from feedback linearization (Appendix D), sliding control (Appendix E) etc.

The design of the adaptation law typically comes from analyzing the dynamics of the tracking error, which as we will see often appears in the form of Lemma 6.1.

The convergence of the closed-loop system is usually analyzed with the help of a Lyapunov-like function introduced in Chapter 4.

Lemma 6.1 (Basic Lemma). *Let two signals $e(t)$ and $\phi(t)$ be related by*

$$e(t) = H(p)[k\phi(t)^T v(t)] \quad (6.1)$$

where $e(t)$ a scalar output signal, $H(p)$ a strictly positive real (SPR) transfer function, k an unknown real number with known sign, $\phi(t) \in \mathbb{R}^m$ a control signal, and $v(t) \in \mathbb{R}^m$ a measurable input signal.

If the control signal $\phi(t)$ satisfies

$$\dot{\phi}(t) = -\text{sgn}(k)\gamma e(t)v(t) \quad (6.2)$$

with $\gamma > 0$ a positive constant, then $e(t)$ and $\phi(t)$ are globally bounded. Moreover, if $v(t)$ is bounded, then

$$\lim_{t \rightarrow \infty} e(t) = 0.$$

Proof. Let the state-space representation of (6.1) be

$$\dot{x} = Ax + b[k\phi^T v], \quad e = c^T x. \quad (6.3)$$

Since $H(p)$ is SPR, it follows from the Kalman-Yakubovich Lemma C.1 that there exist $P, Q \succ 0$ such that

$$A^T P + PA = -Q, \quad Pb = c.$$

Let

$$V(x, \phi) = x^T P x + \frac{|k|}{\gamma} \phi^T \phi,$$

clearly V is positive definite (i.e., $V(0, 0) = 0$, and $V(x, \phi) > 0$ for all $x \neq 0, \phi \neq 0$). The time derivative of V along the trajectory defined by (6.3) with ϕ chosen as in (6.2) is

$$\dot{V} = \frac{\partial V}{\partial x} \dot{x} + \frac{\partial V}{\partial \phi} \dot{\phi} \quad (6.4)$$

$$= x^T (PA + A^T P)x + 2x^T Pb(k\phi^T v) + \frac{2|k|}{\gamma} \phi^T (-\text{sgn}(k)\gamma ev) \quad (6.5)$$

$$= -x^T Qx + 2(x^T c)(k\phi^T v) - 2\phi^T (ev) \quad (6.6)$$

$$= -x^T Qx \leq 0. \quad (6.7)$$

As a result, we know x and ϕ must be bounded ($V(x(t), \phi(t)) \leq V(x(0), \phi(0))$ is bounded). Since $e = c^T x$, we know e must be bounded as well.

If the input signal v is also bounded, then \dot{x} is bounded as seen from (6.3). Because $\dot{V} = -2x^T Q\dot{x}$ is now bounded, we know \dot{V} is uniformly continuous. Therefore, by Barbalat's stability certificate (Theorem 4.1), we know \dot{V} tends to zero as t tends to infinity, which implies $\lim_{t \rightarrow \infty} x(t) = 0$ and hence $\lim_{t \rightarrow \infty} e(t) = 0$. \square

6.1.1 First-Order Systems

Consider the first-order single-input single-output (SISO) system

$$\dot{x} = -ax + bu \quad (6.8)$$

where a and b are unknown groundtruth parameters. However, we do assume that the sign of b is known. What if the sign of b is unknown too?

Let $r(t)$ be a reference trajectory, e.g., a step function or a sinusoidal function, and $x_d(t)$ be a desired system trajectory that tracks the reference

$$\dot{x}_d = -a_d x_d + b_d r(t), \quad (6.9)$$

where $a_d, b_d > 0$ are user-defined constants. Note that the transfer function from r to x_d is

$$x_d = \frac{b_d}{p + a_d} r$$

and the system is stable. Review basics of transfer function.

The goal of adaptive control is to design a control law and an adaptation law such that the tracking error of the system $x(t) - x_d(t)$ converges to zero.

Control law. We design the control law as

$$u = \hat{a}_r(t)r + \hat{a}_x(t)x \quad (6.10)$$

where $\hat{a}_r(t)$ and $\hat{a}_x(t)$ are time-varying feedback gains that we wish to adapt. The closed-loop dynamics of system (6.8) with the controller (6.10) is

$$\dot{x} = -ax + b(\hat{a}_r r + \hat{a}_x x) = -(a - b\hat{a}_x)x + b\hat{a}_r r.$$

With the equation above, the reason for choosing the control law (6.10) is clear: if the system parameters (a, b) were known, then choosing

$$a_r^* = \frac{b_d}{b}, \quad a_x^* = \frac{a - a_d}{b} \quad (6.11)$$

leads to the closed-loop dynamics $\dot{x} = -a_d x + b_d r$ that is exactly what we want in (6.9).

However, in adaptive control, since the true parameters (a, b) are not revealed to the control designer, an adaptation law is needed to dynamically adjust the gains \hat{a}_r and \hat{a}_x based on the tracking error $x(t) - x_d(t)$.

Adaptation law. Let $e(t) = x(t) - x_d(t)$ be the tracking error, and we develop its time derivative

$$\dot{e} = \dot{x} - \dot{x}_d \quad (6.12)$$

$$= -a_d(x - x_d) + (a_d - a + b\hat{a}_x)x + (b\hat{a}_r - b_d)r \quad (6.13)$$

$$= -a_d e + b \underbrace{(\hat{a}_x - \hat{a}_x^*)}_{=: \tilde{a}_x} x + b \underbrace{(\hat{a}_r - \hat{a}_r^*)}_{=: \tilde{a}_r} r \quad (6.14)$$

$$= -a_d e + b(\tilde{a}_x x + \tilde{a}_r r) \quad (6.15)$$

where \tilde{a}_x and \tilde{a}_r are the gain errors w.r.t. the optimal gains in (6.11) if the true parameters were known. The error dynamics (6.15) is equivalent to the following transfer function

$$e = \frac{1}{p + a_d} b(\tilde{a}_x x + \tilde{a}_r r) = \frac{1}{p + a_d} \left(b \begin{bmatrix} \tilde{a}_x \\ \tilde{a}_r \end{bmatrix}^T \begin{bmatrix} x \\ r \end{bmatrix} \right), \quad (6.16)$$

which is in the form of (6.1). Therefore, we choose the adaptation law

$$\begin{bmatrix} \dot{\tilde{a}_x} \\ \dot{\tilde{a}_r} \end{bmatrix} = -\text{sgn}(b)\gamma e \begin{bmatrix} x \\ r \end{bmatrix}. \quad (6.17)$$

Tracking convergence. With the control law (6.10) and the adaptation law (6.17), we can prove that the tracking error converges to zero, using Lemma 6.1. With $\tilde{a} = [\tilde{a}_x, \tilde{a}_r]^T$, let

$$V(e, \tilde{a}) = e^2 + \frac{|b|}{\gamma} \tilde{a}^T \tilde{a} \quad (6.18)$$

be a positive definite Lyapunov function candidate with time derivative

$$\dot{V} = -2a_d e^2 \leq 0.$$

Clearly, e and \tilde{a} are both bounded. Assuming the reference trajectory r is bounded, we know x_d is bounded (due to (6.9)) and hence x is bounded (due to $e = x - x_d$ being bounded). Consequently, from the error dynamics (6.15) we know \dot{e} is bounded, which implies $\dot{V} = -4a_d e \dot{e}$ is bounded and \dot{V} is uniformly continuous. By Barbalat's stability certificate 4.1, we conclude $e(t) \rightarrow 0$ as $t \rightarrow \infty$.

It is always better to combine mathematical analysis with intuitive understanding. Can you explain intuitively why the adaptation law (6.17) makes sense? (Hint: think about how the control should react to a negative/positive tracking error.)

Parameter convergence. We have shown the control law (6.10) and the adaptation law (6.17) guarantee to track the reference trajectory. However, is it guaranteed that the gains of the controller (6.10) also converge to the optimal gains in (6.11)?

We will now show that the answer is indefinite and it depends on the reference trajectory $r(t)$. Because the tracking error e converges to zero, and e is the output of a stable filter (6.16), we know the input $b(\tilde{a}_x x + \tilde{a}_r r)$ must also converge to zero. On the other hand, the adaptation law (6.17) shows that both $\dot{\tilde{a}}_x$ and $\dot{\tilde{a}}_r$ converge to zero (due to e converging to zero and x, r being bounded). As a result, we know $\tilde{a} = [\tilde{a}_x, \tilde{a}_r]^T$ converges to a constant that satisfies

$$v^T \tilde{a} = 0, \quad v = \begin{bmatrix} x \\ r \end{bmatrix}, \quad (6.19)$$

which is a single linear equation of \tilde{a} with time-varying coefficients.

- **Constant reference: no guaranteed convergence.** Suppose $r(t) \equiv r_0 \neq 0$ for all t . From (6.9) we know $x = x_d = \alpha r_0$ when $t \rightarrow \infty$, where α is the constant DC gain of the stable filter. Therefore, the linear equation (6.19) reduces to

$$\alpha \tilde{a}_x + \tilde{a}_r = 0.$$

This implies that \tilde{a} does not necessarily converge to zero. In fact, it converges to a straight line in the parameter space.

- **Persistent excitation: guaranteed convergence.** However, when the signal v satisfies the so-called *persistent excitation* condition, which states that for any t , there exists $T, \beta > 0$ such that

$$\int_t^{t+T} v v^T d\tau \geq \beta I, \quad (6.20)$$

then \tilde{a} is guaranteed to converge to zero. To see this, we multiply (6.19) by v and integrate it from t to $t + T$, which gives rise to

$$\left(\int_t^{t+T} vv^T d\tau \right) \tilde{a} = 0.$$

By the persistent excitation condition (6.20), we infer that $\tilde{a} = 0$ is the only solution.

It remains to understand under what conditions of the reference trajectory $r(t)$ can we guarantee the persistent excitation of v . We leave it as an exercise for the reader to show, if $r(t)$ contains at least one sinusoidal component, then the persistent excitation condition of v is guaranteed.

Exercise 6.1 (Extension to Nonlinear Systems). Design a control law and an adaptation law for the following system

$$\dot{x} = -ax - cf(x) + bu$$

with unknown true parameters (a, b, c) (assume the sign of b is known) and known nonlinearity $f(x)$ to track a reference trajectory $r(t)$. Analyze the convergence of tracking error and parameter estimation error.

6.1.2 High-Order Systems

Consider an n -th order nonlinear system

$$q^{(n)} + \sum_{i=1}^n \alpha_i f_i(x, t) = bu \quad (6.21)$$

where $x = [q, \dot{q}, \ddot{q}, \dots, q^{(n-1)}]^T$ is the state of the system, f_i 's are known nonlinearities, $(\alpha_1, \dots, \alpha_n, b)$ are unknown parameters of the system (with $\text{sgn}(b)$ known).

The goal of adaptive control is to control the system (6.21) trajectory to follow a desired trajectory $q_d(t)$ despite not knowing the true parameters.

To facilitate the derivation of the adaptive controller, let us divide both sides of (6.21) by b

$$hq^{(n)} + \sum_{i=1}^n a_i f_i(x, t) = u \quad (6.22)$$

where $h = 1/b$ and $a_i = \alpha_i/b$.

Control law. Recall that the choice of the control law is typically inspired by the control design if the true system parameters were known. We will borrow ideas from sliding control (Appendix E).

- **Known parameters.** Let $e = q(t) - q_d(t)$ be the tracking error, and define the following combined error

$$s = e^{(n-1)} + \lambda_{n-2}e^{(n-2)} + \cdots + \lambda_0 e = \Delta(p)e$$

where $\Delta(p) = p^{n-1} + \lambda_{n-2}p^{(n-2)} + \cdots + \lambda_0$ is a stable polynomial with user-chosen coefficients $\lambda_0, \dots, \lambda_{n-2}$. The rationale for defining the combined error s is that the convergence of e to zero can be guaranteed by the convergence of s to zero (when $\Delta(p)$ is stable). Note that s can be equivalently written as

$$s = (q^{(n-1)} - q_d^{(n-1)}) + \lambda_{n-2}e^{(n-2)} + \cdots + \lambda_0 e \quad (6.23)$$

$$= q^{(n-1)} - \underbrace{\left(q_d^{(n-1)} - \lambda_{n-2}e^{(n-2)} - \cdots - \lambda_0 e \right)}_{q_r^{(n-1)}}. \quad (6.24)$$

Now consider the control law

$$u = hq_r^{(n)} - ks + \sum_{i=1}^n a_i f_i(x, t) \quad (6.25)$$

where

$$q_r^{(n)} = q_d^{(n)} - \lambda_{n-2}e^{(n-1)} - \cdots - \lambda_0 \dot{e}$$

and k is a design constant that has the same sign as h . This choice of control, plugged into the system dynamics (6.22), leads to

$$hq^{(n)} + \sum_{i=1}^n a_i f_i(x, t) = hq_r^{(n)} - ks + \sum_{i=1}^n a_i f_i(x, t) \iff \quad (6.26)$$

$$h \left(q^{(n)} - q_r^{(n)} \right) + ks = 0 \iff \quad (6.27)$$

$$h\dot{s} + ks = 0, \quad (6.28)$$

which guarantees the exponential convergence of s to zero (note that h and k have the same sign), and hence the convergence of e to zero.

- **Unknown parameters.** Inspired by the control law with known parameters in (6.25), we design the adaptive control law as

$$u = \hat{h}q_r^{(n)} - ks + \sum_{i=1}^n \hat{a}_i f_i(x, t), \quad (6.29)$$

where the time-varying gains $\hat{h}, \hat{a}_1, \dots, \hat{a}_n$ will be adjusted by an adaptation law.

Adaptation law. Inserting the adaptive control law (6.29) into the system dynamics (6.22), we obtain

$$h\dot{s} + ks = \tilde{h}q_r^{(n)} + \sum_{i=1}^n \tilde{a}_i f_i(x, t) \iff \quad (6.30)$$

$$s = \frac{1}{p + k/h} \frac{1}{h} \underbrace{\begin{pmatrix} \tilde{h} \\ \tilde{a}_1 \\ \vdots \\ \tilde{a}_n \end{pmatrix}^T \begin{pmatrix} q_r^{(n)} \\ f_1(x, t) \\ \vdots \\ f_n(x, t) \end{pmatrix}}_{=: \phi^T v} \quad (6.31)$$

where $\tilde{h} = \hat{h} - h$ and $\tilde{a}_i = \hat{a}_i - a_i, i = 1, \dots, n$. Again, (6.31) is in the familiar form of (6.1), which naturally leads to the following adaptation law with $\gamma > 0$ a chosen constant

$$\dot{\phi} = \begin{bmatrix} \dot{\tilde{h}} \\ \dot{\tilde{a}}_1 \\ \vdots \\ \dot{\tilde{a}}_n \end{bmatrix} = -\text{sgn}(h)\gamma s \begin{bmatrix} q_r^{(n)} \\ f_1(x, t) \\ \vdots \\ f_n(x, t) \end{bmatrix}. \quad (6.32)$$

Tracking and parameter convergence. With the following Lyapunov function

$$V(s, \phi) = |h|s^2 + \frac{1}{\gamma} \phi^T \phi, \quad \dot{V}(s, \phi) = -2|k|s^2, \quad (6.33)$$

the global convergence of s to zero can be easily shown. For parameter convergence, it is easy to see that when v satisfies the persistent excitation condition, we have that ϕ converges to zero. (However, the relationship between the reference trajectory $q_d(t)$ and the persistent excitation of v becomes nontrivial due to the nonlinearities f_i .)

6.1.3 Robotic Manipulator

So far our focus has been on systems with a single input ($u \in \mathbb{R}$). In the following, we will show that similar techniques can be applied to adaptive control of systems with multiple inputs, particularly, trajectory control of a robotic manipulator.

Let $q \in \mathbb{R}^n$ be the joint angles of a multi-link robotic arm, and $\dot{q} \in \mathbb{R}^n$ be the joint velocities. The dynamics of a robotic manipulator reads

$$H(q)\ddot{q} + C(q, \dot{q})\dot{q} + g(q) = \tau, \quad (6.34)$$

where $H(q) \in \mathbb{S}_{++}^n$ is the manipulator inertia matrix (that is positive definite), $C(q, \dot{q})\dot{q}$ is a vector of centripetal and Coriolis torques (with $C(q, \dot{q}) \in \mathbb{R}^{n \times n}$), and $g(q)$ denotes gravitational torques.




Figure 6.1: Planar two-link manipulator

Example 6.1 (Planar Two-link Manipulator). The dynamics of a planar two-link manipulator in Fig. 6.1 is

$$\begin{bmatrix} H_{11} & H_{12} \\ H_{21} & H_{22} \end{bmatrix} \begin{bmatrix} \ddot{q}_1 \\ \ddot{q}_2 \end{bmatrix} + \begin{bmatrix} -h\dot{q}_2 & -h(\dot{q}_1 + \dot{q}_2) \\ h\dot{q}_1 & 0 \end{bmatrix} \begin{bmatrix} \dot{q}_1 \\ \dot{q}_2 \end{bmatrix} = \begin{bmatrix} \tau_1 \\ \tau_2 \end{bmatrix}, \quad (6.35)$$

where

$$H_{11} = a_1 + 2a_3 \cos q_2 + 2a_4 \sin q_2 \quad (6.36)$$

$$H_{12} = H_{21} = a_2 + a_3 \cos q_2 + a_4 \sin q_2 \quad (6.37)$$

$$H_{22} = a_2 \quad (6.38)$$

$$h = a_3 \sin q_2 - a_4 \cos q_2 \quad (6.39)$$

with

$$a_1 = I_1 + m_1 l_{c1}^2 + I_e + m_e l_{ce}^2 + m_e l_1^2 \quad (6.40)$$

$$a_2 = I_e + m_e l_{ce}^2 \quad (6.41)$$

$$a_3 = m_e l_1 l_{ce} \cos \delta_e \quad (6.42)$$

$$a_4 = m_e l_1 l_{ce} \sin \delta_e. \quad (6.43)$$

As seen from the above example, the parameters a (which are nonlinear functions of the physical parameters such as mass and length) enter linearly in H and C ($g(q)$ is ignored because the manipulator is on a horizontal plane).

The goal of the control design is to have the manipulator track a desired trajectory $q_d(t)$.

Known parameters. When the parameters are known, we follow the sliding control design framework. Let $\tilde{q} = q(t) - q_d(t)$ be the tracking error, and define the combined error

$$s = \dot{\tilde{q}} + \Lambda \tilde{q} = \dot{q} - \underbrace{(\dot{q}_d - \Lambda \tilde{q})}_{\dot{q}_r}$$

where $\Lambda \in \mathbb{S}_{++}^n$ is a user-chosen positive definite matrix (in general we want $-\Lambda$ to be Hurwitz). In this case, $s \rightarrow 0$ implies $\tilde{q} \rightarrow 0$ as $t \rightarrow \infty$. Choosing the

control law (coming from feedback linearization Appendix D)

$$\tau = H\ddot{q}_r - K_D s + C\dot{q} + g(q) \quad (6.44)$$

with $K_D \in \mathbb{S}_{++}^n$ positive definite leads to the closed-loop dynamics

$$H\dot{s} + K_D s = 0 \iff \dot{s} = -H^{-1}K_D s.$$

Because the matrix $H^{-1}K_D$ is the product of two positive definite matrices (recall H is positive definite and so is H^{-1}), it has strictly positive real eigenvalues.¹ Hence, $-H^{-1}K_D$ is Hurwitz and s is guaranteed to converge to zero.

Control law. A closer look at the controller (6.44) allows us to write it in the following form

$$\tau = H\ddot{q}_r + C(s + \dot{q}_r) + g(q) - K_D s \quad (6.45)$$

$$= H\ddot{q}_r + C\dot{q}_r + g(q) + (C - K_D)s \quad (6.46)$$

$$= Y(q, \dot{q}, \ddot{q}_r)a + (C - K_D)s \quad (6.47)$$

where $a \in \mathbb{R}^m$ contains all the parameters and $Y \in \mathbb{R}^{n \times m}$ is the matrix that collects all the coefficients of a in $H\ddot{q}_r + C\dot{q}_r + g(q)$. As a result, we design the adaptive control law to be

$$\tau = Y\hat{a} - K_D s, \quad (6.48)$$

with \hat{a} the time-varying parameter that we wish to adapt. Note that here we have done something strange: the adaptive control law does not exactly follow the controller (6.44) in the known-parameter case.² We first separated s from \dot{q} and wrote $Y a = H\ddot{q}_r + C\dot{q}_r + g$ instead of $Y a = H\ddot{q}_r + C\dot{q} + g$; then we dropped the “ C ” matrix in front of s in the adaptive control law. The reason for doing this will soon become clear when we analyze the tracking convergence.

Adaptation law and tracking convergence. Recall that the key of adaptive control is to design a control law and an adaptation law such that global convergence of the tracking error s can be guaranteed by a Lyapunov function. Looking at the previous Lyapunov functions in (6.18) and (6.33), we see that they both contain a positive definite term in the tracking error s (or e if in first-order systems) and another positive definite term in the parameter error \tilde{a} . This hints us that we may try a Lyapunov candidate function of the following form

$$V = \frac{1}{2} (s^T H s + \tilde{a}^T \Gamma^{-1} \tilde{a}), \quad (6.49)$$

¹Consider two positive definite matrices A and B , let $B = B^{1/2} B^{1/2}$. The product AB can be written as $AB = AB^{1/2} B^{1/2} = B^{-1/2} (B^{1/2} A B^{1/2}) B^{1/2}$. Therefore AB is similar to $B^{1/2} A B^{1/2}$ and is positive definite.

²In fact, one can show that the controller (6.48) with known parameters, i.e., $\tau = Y a - K_D s$, also guarantees the convergence of s towards zero, though it is different from the feedback linearization controller (6.44). Try proving the convergence with a Lyapunov candidate $V = \frac{1}{2} s^T H s$.

where $\Gamma \in \mathbb{S}_{++}^m$ is a constant positive definite matrix, and $\tilde{a} = \hat{a} - a$ is the parameter error.

The next step would be to derive the time derivative of V , which, as we can expect, will contain a term that involves \dot{H} and complicates our analysis. Fortunately, the following lemma will help us.

Lemma 6.2. *For the manipulator dynamics (6.34), there exists a way to define C such that $\dot{H} - 2C$ is skew-symmetric.*

Proof. See Section 9.1, page 399-402 in (?). You should also check if this is true for the planar two-link manipulator dynamics in Example 6.1. \square

With Lemma 6.2, the time derivative of V in (6.49) reads

$$\dot{V} = s^T H \dot{s} + \frac{1}{2} s^T \dot{H} s + \tilde{a}^T \Gamma^{-1} \dot{\tilde{a}} \quad (6.50)$$

$$= s^T (H \ddot{q} - H \ddot{q}_r) + \frac{1}{2} s^T \dot{H} s + \tilde{a}^T \Gamma^{-1} \dot{\tilde{a}} \quad (6.51)$$

$$= s^T (\tau - C \dot{q} - g - H \ddot{q}_r) + \frac{1}{2} s^T \dot{H} s + \tilde{a}^T \Gamma^{-1} \dot{\tilde{a}} \quad (6.52)$$

$$= s^T (\tau - H \ddot{q}_r - C(s + \dot{q}_r) - g) + \frac{1}{2} s^T \dot{H} s + \tilde{a}^T \Gamma^{-1} \dot{\tilde{a}} \quad (6.53)$$

$$= s^T (\tau - H \ddot{q}_r - C \dot{q}_r - g) + \frac{1}{2} s^T (\dot{H} - 2C) s + \tilde{a}^T \Gamma^{-1} \dot{\tilde{a}} \quad (6.54)$$

$$= s^T (\tau - H \ddot{q}_r - C \dot{q}_r - g) + \tilde{a}^T \Gamma^{-1} \dot{\tilde{a}} \quad (6.55)$$

$$= s^T (Y \hat{a} - K_D s - Y a) + \tilde{a}^T \Gamma^{-1} \dot{\tilde{a}} \quad (6.56)$$

$$= s^T Y \tilde{a} + \tilde{a}^T \Gamma^{-1} \dot{\tilde{a}} - s^T K_D s, \quad (6.57)$$

where we used the manipulator dynamics (6.34) to rewrite $H \ddot{q}$ in (6.52), used $\dot{H} - 2C$ is skew-symmetric in (6.54), invoked the adaptive control law (6.48) and reused $Y a = H \ddot{q}_r + C \dot{q}_r + g(q)$ in (6.56). The derivation above explains why the choice of the control law in (6.48) did not exactly follow its counterpart when the parameters are known: we need to use $s^T C s$ to cancel $\frac{1}{2} s^T \dot{H} s$ in (6.54).

We then wonder if we can design $\dot{\tilde{a}}$ such that \dot{V} in (6.57) is negative semidefinite? This turns out to be straightforward with the adaptation law

$$\dot{\tilde{a}} = -\Gamma Y^T s, \quad (6.58)$$

to make $s^T Y \tilde{a} + \tilde{a}^T \Gamma^{-1} \dot{\tilde{a}}$ vanish and so

$$\dot{V} = -s^T K_D s \leq 0.$$

We are not done yet. To show s converges to zero (which is implied by \dot{V} converges to zero), by Barbalat's stability certificate 4.1, it suffices to show

$$\ddot{V} = -2s^T K_D \dot{s}$$

is bounded. We already know s and \tilde{a} are bounded, due to the fact that V in (6.49) is bounded. Therefore, we only need to show \dot{s} is bounded. To do so, we plug the adaptive control law (6.48) into the manipulator dynamics (6.34) and obtain

$$H\dot{s} + (C + K_D)s = Y\tilde{a},$$

which implies the boundedness of \dot{s} (note that H is uniformly positive definite, i.e., $H \succeq \alpha I$ for some $\alpha > 0$). This concludes the analysis of the tracking convergence $s \rightarrow 0$ as $t \rightarrow \infty$.

6.2 Certainty-Equivalent Adaptive Control

Appendix A

Linear System Theory

Consider the linear time-invariant (LTI) system

$$\dot{x} = Ax + Bu, \quad y = Cx + Du, \quad (\text{A.1})$$

where $x \in \mathbb{R}^n$ the state, $u \in \mathbb{R}^m$ the control, and A, B, C, D are constant matrices with proper sizes.

Theorem A.1 (Detectability). *The following are equivalent for the LTI system (A.1):*

1. (A, C) is detectable;
2. The matrix $\begin{bmatrix} A - \lambda I \\ C \end{bmatrix}$ has full column rank for all $\lambda \in \mathbb{C}$ such that $\text{Re}(\lambda) \geq 0$;
;
3. For all λ and x such that $Ax = \lambda x$ and $\text{Re}(\lambda) \geq 0$, $Cx \neq 0$;
4. There exists a matrix L such that $A + LC$ is Hurwitz;
5. (A^*, C^*) is stabilizable.

Appendix B

Convex Analysis and Optimization

Appendix C

The Kalman-Yakubovich Lemma

Lemma C.1 (Kalman-Yakubovich). *Consider a controllable linear time-invariant system*

$$\dot{x} = Ax + by = c^T x.$$

The transfer function

$$h(p) = c^T (pI - A)^{-1} b$$

is strictly positive real (SPR) if and only if there exist positive definite matrices P and Q such that

$$A^T P + PA = -QPb = c.$$

Appendix D

Feedback Linearization

Appendix E

Sliding Control